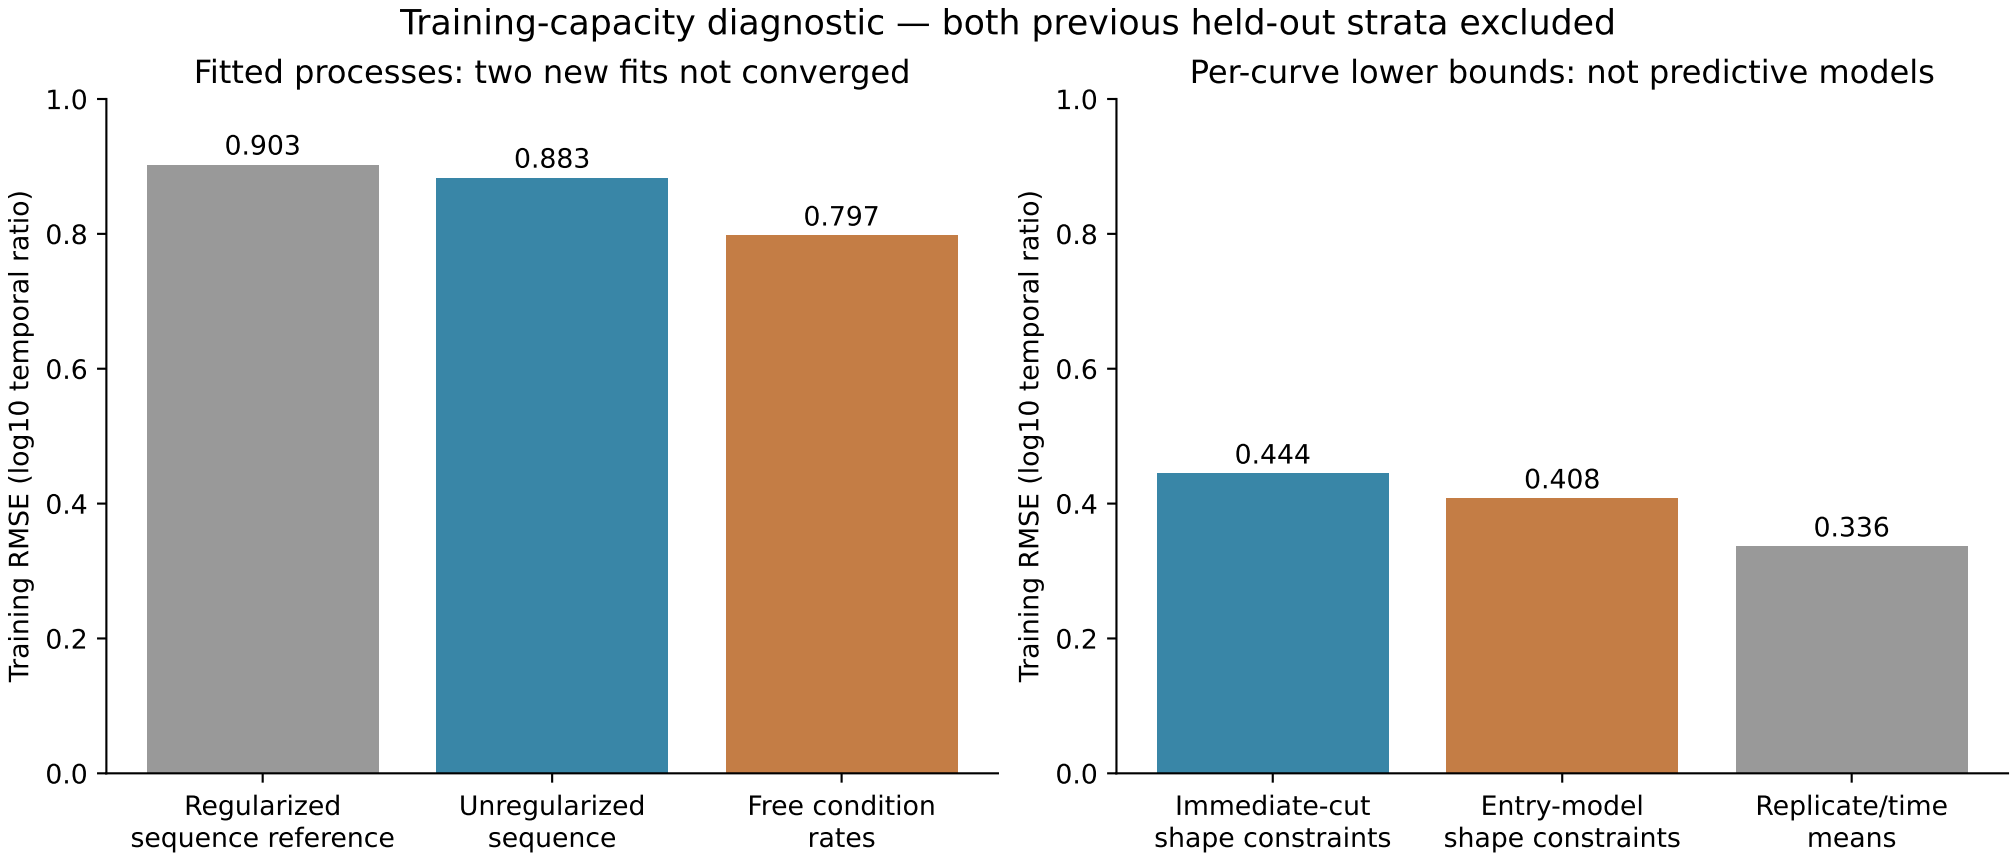
Training-capacity diagnostic — both previous held-out strata excluded
Fitted processes: two new fits not converged
Per-curve lower bounds: not predictive models
Training RMSE (log10 temporal ratio)
1.0
0.8
0.6
0.4
0.2
0.0
0.903
0.883
0.797
0.444
0.408
0.336
Regularized sequence reference
Unregularized sequence
Free condition rates
Immediate-cut shape constraints
Entry-model shape constraints
Replicate/time means